

# Data Preparation

In [1]: `!wget https://raw.githubusercontent.com/keerthy456/Machine-Learning-Final-Project-Vakkalagadda-Keerthi/main/heart_disease.csv`

```
--2022-05-06 04:10:42-- https://raw.githubusercontent.com/keerthy456/Machine-Learning-Final-Project-Vakkalagadda-Keerthi/main/heart_disease.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.111.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.111.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 25189554 (24M) [text/plain]
Saving to: 'heart_disease.csv.2'
```

```
heart_disease.csv.2 100%[=====>] 24.02M 118MB/s in 0.2s
```

```
2022-05-06 04:10:44 (118 MB/s) - 'heart_disease.csv.2' saved [25189554/25189554]
```



In [3]: `import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split`


```
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
from sklearn.metrics import confusion_matrix

from sklearn.linear_model import LinearRegression ,LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

In [4]: `heart_df = pd.read_csv('heart_disease.csv')`

```
In [5]: plt.figure(figsize = (5,5))  
sns.countplot(x = heart_df['HeartDisease'], hue = 'HeartDisease', data = heart_df)
```

```
Out[5]: <AxesSubplot:xlabel='HeartDisease', ylabel='count'>
```



### Sampling Dataset

```
In [6]: class_no = heart_df[heart_df['HeartDisease'] == 'No']  
class_yes = heart_df[heart_df['HeartDisease'] == 'Yes']
```

```
In [130]: len(class_no)
```


```
Out[130]: 292422
```

```
In [7]: class_no = class_no.sample(len(class_yes), replace=False)  
new_df = pd.concat([class_no, class_yes], axis=0)  
print('Target class Distribution after Sampling in :')  
print(new_df['HeartDisease'].value_counts())  
heart_df = new_df.copy()
```

```
Target class Distribution after Sampling in :  
No      27373  
Yes     27373  
Name: HeartDisease, dtype: int64
```

```
In [172]: plt.figure(figsize = (5,5))  
sns.countplot(x = new_df['HeartDisease'], hue = 'HeartDisease', data = new_df)
```


Out[172]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f5f6befde90>



In [210]: `pip install dython`

```
Requirement already satisfied: dython in /usr/local/lib/python3.7/dist-packages (0.7.1.post3)
Requirement already satisfied: pandas>=1.3.2 in /usr/local/lib/python3.7/dist-packages (from dython) (1.3.5)
Requirement already satisfied: scipy>=1.7.1 in /usr/local/lib/python3.7/dist-packages (from dython) (1.7.3)
Requirement already satisfied: matplotlib>=3.4.3 in /usr/local/lib/python3.7/dist-packages (from dython) (3.5.2)
Requirement already satisfied: scikit-plot>=0.3.7 in /usr/local/lib/python3.7/dist-packages (from dython) (0.3.7)
Requirement already satisfied: seaborn>=0.11.0 in /usr/local/lib/python3.7/dist-packages (from dython) (0.11.2)
Requirement already satisfied: scikit-learn>=0.24.2 in /usr/local/lib/python3.7/dist-packages (from dython) (1.0.2)
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.7/dist-packages (from dython) (1.21.6)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (2.8.2)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (0.11.0)
Requirement already satisfied: pyparsing>=2.2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (3.0.8)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (7.1.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (1.4.2)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (4.33.3)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (21.3)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from kiwisolver>=1.0.1->matplotlib>=3.4.3->dython) (4.2.0)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-packages (from pandas>=1.3.2->dython) (2022.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.7->matplotlib>=3.4.3->dython) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.24.2->dython) (3.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.24.2->dython) (1.1.0)
```

```
In [215]: from dython.nominal import associations  
          associations(heart_df, figsize=(15,15))
```



```

Out[215]: {'ax': <matplotlib.axes._subplots.AxesSubplot at 0x7f5f6b3c2c50>,
'corr':
Stroke \
HeartDisease      1.000000  0.086123  0.186801      0.062757  0.22633
1
BMI                0.086123  1.000000  0.011327      0.035059  0.01952
9
Smoking            0.186801  0.011327  1.000000      0.083308  0.07455
1
AlcoholDrinking    0.062757  0.035059  0.083308      1.000000  0.01942
8
Stroke             0.226331  0.019529  0.074551      0.019428  1.00000
0
PhysicalHealth      0.241695  0.119827  0.124253      0.024145  0.17045
4
MentalHealth        0.047635  0.076881  0.083622      0.038517  0.07334
7
DiffWalking         0.287334  0.197973  0.128690      0.042537  0.20580
7
Sex                 0.123744  0.019856  0.111418      0.005145  0.00000
0
AgeCategory         0.440379  0.186255  0.165518      0.090602  0.14229
1
Race                0.103719  0.105677  0.115258      0.028900  0.05316
6
Diabetic            0.264801  0.234813  0.077663      0.073200  0.12789
5
PhysicalActivity     0.163729  0.162181  0.104091      0.015783  0.10712
7
GenHealth           0.398104  0.218810  0.200787      0.042866  0.22529
1
SleepTime           0.014021 -0.043605  0.023065      0.002619  0.00869
8
Asthma              0.072932  0.110901  0.030711      0.009336  0.05915
0
KidneyDisease       0.181635  0.072841  0.034510      0.036570  0.11055
8
SkinCancer          0.143247  0.040863  0.027861      0.010678  0.04682
6

```

	PhysicalHealth	MentalHealth	DiffWalking	Sex \
HeartDisease	0.241695	0.047635	0.287334	0.123744
BMI	0.119827	0.076881	0.197973	0.019856
Smoking	0.124253	0.083622	0.128690	0.111418
AlcoholDrinking	0.024145	0.038517	0.042537	0.005145
Stroke	0.170454	0.073347	0.205807	0.000000
PhysicalHealth	1.000000	0.324227	0.453900	0.042898
MentalHealth	0.324227	1.000000	0.194129	0.110779
DiffWalking	0.453900	0.194129	1.000000	0.081818
Sex	0.042898	0.110779	0.081818	1.000000
AgeCategory	0.143431	0.157536	0.253090	0.057654
Race	0.056797	0.061992	0.075858	0.067554
Diabetic	0.192377	0.054333	0.248460	0.091042
PhysicalActivity	0.276534	0.126524	0.325859	0.063045
GenHealth	0.640061	0.298198	0.496175	0.029681
SleepTime	-0.070978	-0.126860	0.033283	0.005607
Asthma	0.142356	0.123548	0.144099	0.087855

KidneyDisease	0.176063	0.057116	0.190639	0.009989
SkinCancer	0.040202	0.032021	0.063370	0.056208

	AgeCategory	Race	Diabetic	PhysicalActivity	\
HeartDisease	0.440379	0.103719	0.264801	0.163729	
BMI	0.186255	0.105677	0.234813	0.162181	
Smoking	0.165518	0.115258	0.077663	0.104091	
AlcoholDrinking	0.090602	0.028900	0.073200	0.015783	
Stroke	0.142291	0.053166	0.127895	0.107127	
PhysicalHealth	0.143431	0.056797	0.192377	0.276534	
MentalHealth	0.157536	0.061992	0.054333	0.126524	
DiffWalking	0.253090	0.075858	0.248460	0.325859	
Sex	0.057654	0.067554	0.091042	0.063045	
AgeCategory	1.000000	0.106135	0.137154	0.144289	
Race	0.106135	1.000000	0.045507	0.056471	
Diabetic	0.137154	0.045507	1.000000	0.171012	
PhysicalActivity	0.144289	0.056471	0.171012	1.000000	
GenHealth	0.128188	0.056113	0.181264	0.330934	
SleepTime	0.157892	0.053026	0.017061	0.006354	
Asthma	0.065954	0.060598	0.072943	0.069547	
KidneyDisease	0.138081	0.033205	0.192191	0.114242	
SkinCancer	0.280373	0.160454	0.035720	0.000000	

	GenHealth	SleepTime	Asthma	KidneyDisease	SkinCancer
HeartDisease	0.398104	0.014021	0.072932	0.181635	0.143247
BMI	0.218810	-0.043605	0.110901	0.072841	0.040863
Smoking	0.200787	0.023065	0.030711	0.034510	0.027861
AlcoholDrinking	0.042866	0.002619	0.009336	0.036570	0.010678
Stroke	0.225291	0.008698	0.059150	0.110558	0.046826
PhysicalHealth	0.640061	-0.070978	0.142356	0.176063	0.040202
MentalHealth	0.298198	-0.126860	0.123548	0.057116	0.032021
DiffWalking	0.496175	0.033283	0.144099	0.190639	0.063370
Sex	0.029681	0.005607	0.087855	0.009989	0.056208
AgeCategory	0.128188	0.157892	0.065954	0.138081	0.280373
Race	0.056113	0.053026	0.060598	0.033205	0.160454
Diabetic	0.181264	0.017061	0.072943	0.192191	0.035720
PhysicalActivity	0.330934	0.006354	0.069547	0.114242	0.000000
GenHealth	1.000000	0.064136	0.173484	0.232674	0.059565
SleepTime	0.064136	1.000000	0.054742	0.000934	0.039736
Asthma	0.173484	0.054742	1.000000	0.059647	0.000000
KidneyDisease	0.232674	0.000934	0.059647	1.000000	0.067498
SkinCancer	0.059565	0.039736	0.000000	0.067498	1.000000

By looking at the above coorelation matrix I believe 'Alcohol Drinking', 'Mental Health', 'Sleep Time', 'Race' doesnot seem to be highly coorelated (individually/when combined with other features as well) with 'Target Feature - Heart Disease'. So I'll be dropping them from my dataset and train the model with remaining features.

```
In [8]: numeric_features = heart_df.select_dtypes(include=[np.number])
```

```
In [134]: categorical_features= [col for col in heart_df.columns if heart_df[col].dtypes
== 'object']
```



In [124]: `heart_df.isnull().any()`

```
Out[124]: HeartDisease      False
BMI                  False
Smoking              False
AlcoholDrinking      False
Stroke               False
PhysicalHealth        False
MentalHealth          False
DiffWalking          False
Sex                  False
AgeCategory           False
Race                  False
Diabetic              False
PhysicalActivity       False
GenHealth             False
SleepTime             False
Asthma                False
KidneyDisease         False
SkinCancer            False
dtype: bool
```

In [9]: `heart_df = heart_df.drop(columns=['AlcoholDrinking', 'SleepTime', 'MentalHealth', 'Race'])`

In [10]: `heart_df['heartdisease_GenHealth'] = heart_df.groupby('GenHealth')['HeartDisease'].transform('count')`  
`heart_df['mean_PhysicalHealth'] = heart_df.groupby('DiffWalking')['PhysicalHealth'].transform('mean')`  
`heart_df['BMI_Std'] = heart_df.groupby('PhysicalActivity')['BMI'].transform('std')`

In [11]: `heart_df.groupby('GenHealth').count()`

Out[11]:

	HeartDisease	BMI	Smoking	Stroke	PhysicalHealth	DiffWalking	Sex	AgeCategory
GenHealth								
Excellent	7664	7664	7664	7664	7664	7664	7664	7
Fair	9716	9716	9716	9716	9716	9716	9716	9
Good	17304	17304	17304	17304	17304	17304	17304	17
Poor	4537	4537	4537	4537	4537	4537	4537	4
Very good	15525	15525	15525	15525	15525	15525	15525	15

In [12]: heart\_df.head()

Out[12]:

	HeartDisease	BMI	Smoking	Stroke	PhysicalHealth	DiffWalking	Sex	AgeCategory
249975	No	24.96	No	No	0.0	No	Female	65-69
311944	No	28.19	Yes	No	0.0	No	Female	65-69
63414	No	23.40	Yes	No	0.0	No	Male	70-74
175588	No	30.52	Yes	No	0.0	No	Male	70-74
144586	No	28.97	Yes	No	0.0	No	Male	50-54

In [13]: `print('\nCategorical Columns\n')`  
`heart_df.select_dtypes(include=['O']).nunique()`

Categorical Columns

Out[13]:

HeartDisease	2
Smoking	2
Stroke	2
DiffWalking	2
Sex	2
AgeCategory	13
Diabetic	4
PhysicalActivity	2
GenHealth	5
Asthma	2
KidneyDisease	2
SkinCancer	2
dtype: int64	

From the above output we can see some of the categorical features are binary class variables which has 2 unique values -yes/no, and some have more than 2 class values. So, for encoding binary features I've used Label encoder for converting yes/no to either 1 or 0 and dummies technique is used for categorical columns which has more than two unique values. Coming to dropping the unnecessary columns, I did not drop any columns because the dataset does not have any unique identifiers/ patient name/ address/zip code information and based on my intuition everything seems important. However the age Category is in categorical(each value is specified as range) and I'm converting it into integer by taking the mean if the given range.

In [14]: `for col in ['HeartDisease', 'Smoking', 'Stroke', 'DiffWalking', 'Sex', 'PhysicalActivity', 'Asthma', 'KidneyDisease', 'SkinCancer']:`  
 `if heart_df[col].dtype == 'O':`  
 `le = LabelEncoder()`  
 `heart_df[col] = le.fit_transform(heart_df[col])`

In [15]: heart\_df.head()

Out[15]:

	HeartDisease	BMI	Smoking	Stroke	PhysicalHealth	DiffWalking	Sex	AgeCategory	I
249975	0	24.96	0	0	0.0	0	0	65-69	
311944	0	28.19	1	0	0.0	0	0	65-69	
63414	0	23.40	1	0	0.0	0	1	70-74	
175588	0	30.52	1	0	0.0	0	1	70-74	
144586	0	28.97	1	0	0.0	0	1	50-54	

In [16]: categoricals = heart\_df[['GenHealth', 'Diabetic']]  
categoricals.head()

Out[16]:

	GenHealth	Diabetic
249975	Good	No
311944	Excellent	No
63414	Very good	No
175588	Good	Yes
144586	Good	No

In [17]: cat\_dummies = pd.get\_dummies(categoricals, drop\_first=True)  
cat\_dummies.head()

Out[17]:


	GenHealth_Fair	GenHealth_Good	GenHealth_Poor	GenHealth_Very good	Diabetic_No, borderline diabetes	Diabet
249975	0	1	0	0	0	
311944	0	0	0	0	0	
63414	0	0	0	1	0	
175588	0	1	0	0	0	
144586	0	1	0	0	0	

```
In [18]: # Drop the redundant columns
heart_df.drop(list(categoricals.columns), axis=1, inplace=True)
# concat the heart and dummies data frames.
heart_df = pd.concat([heart_df, cat_dummies], axis=1)
heart_df.head()
```

Out[18]:

	HeartDisease	BMI	Smoking	Stroke	PhysicalHealth	DiffWalking	Sex	AgeCategory	I
249975	0	24.96	0	0	0.0	0	0	65-69	
311944	0	28.19	1	0	0.0	0	0	65-69	
63414	0	23.40	1	0	0.0	0	1	70-74	
175588	0	30.52	1	0	0.0	0	1	70-74	
144586	0	28.97	1	0	0.0	0	1	50-54	

5 rows × 22 columns




```
In [19]: encode_AgeCategory = {'55-59':57, '80 or older':80, '65-69':67,
                              '75-79':77, '40-44':42, '70-74':72, '60-64':62,
                              '50-54':52, '45-49':47, '18-24':21, '35-39':37,
                              '30-34':32, '25-29':27}
heart_df['AgeCategory'] = heart_df['AgeCategory'].apply(lambda x: encode_AgeCategory[x])
heart_df['AgeCategory'] = heart_df['AgeCategory'].astype('int64')
```

In [20]: heart\_df.head()

Out[20]:

	HeartDisease	BMI	Smoking	Stroke	PhysicalHealth	DiffWalking	Sex	AgeCategory	I
249975	0	24.96	0	0	0.0	0	0	67	
311944	0	28.19	1	0	0.0	0	0	67	
63414	0	23.40	1	0	0.0	0	1	72	
175588	0	30.52	1	0	0.0	0	1	72	
144586	0	28.97	1	0	0.0	0	1	52	

5 rows × 22 columns



## Normalization/Scaling

The range of continuous features are different. Here, I am scaling them to be in-between 0 to 1 by dividing by the maximum value of the respective column

```
In [21]: from sklearn.preprocessing import StandardScaler
num_cols = ['BMI', 'PhysicalHealth']
scaler = StandardScaler()
heart_df[num_cols] = scaler.fit_transform(heart_df[num_cols])
```

```
In [22]: heart_df.describe()[1:][['BMI', 'PhysicalHealth']].T.style.background_gradient(
cmap='Blues')
```

Out[22]:

	mean	std	min	25%	50%	75%	max
<b>BMI</b>	0.000000	1.000009	-2.573955	-0.682142	-0.162125	0.495226	8.948204
<b>PhysicalHealth</b>	-0.000000	1.000009	-0.538968	-0.538968	-0.538968	-0.036554	2.475512

```
In [23]: heart_df = heart_df.drop(columns = ['heartdisease_GenHealth', 'mean_PhysicalHe
alth', 'BMI_Std',])
```

```
In [24]: #Select Features
features = heart_df.drop(columns =['HeartDisease'], axis = 1)
#Select Target
target = heart_df['HeartDisease']
# Set Training and Testing Data
from sklearn.model_selection import train_test_split
```

```
In [25]: high = len(heart_df['HeartDisease']) - sum(heart_df['HeartDisease'])
print("Baseline accuracy : ",high/len(heart_df['HeartDisease']))

Baseline accuracy : 0.5
```

```
In [26]: from sklearn import metrics
```

## Linear Regression

```
In [62]: X_train, X_test, y_train, y_test = train_test_split(features, target, test_siz
e = 0.1, random_state=33)
```

```
In [61]: # plot impact of logloss for single forecasts
from sklearn.metrics import log_loss
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
lin_reg_pred = lin_reg.predict(X_test)
score = lin_reg.score(X_test,y_test)

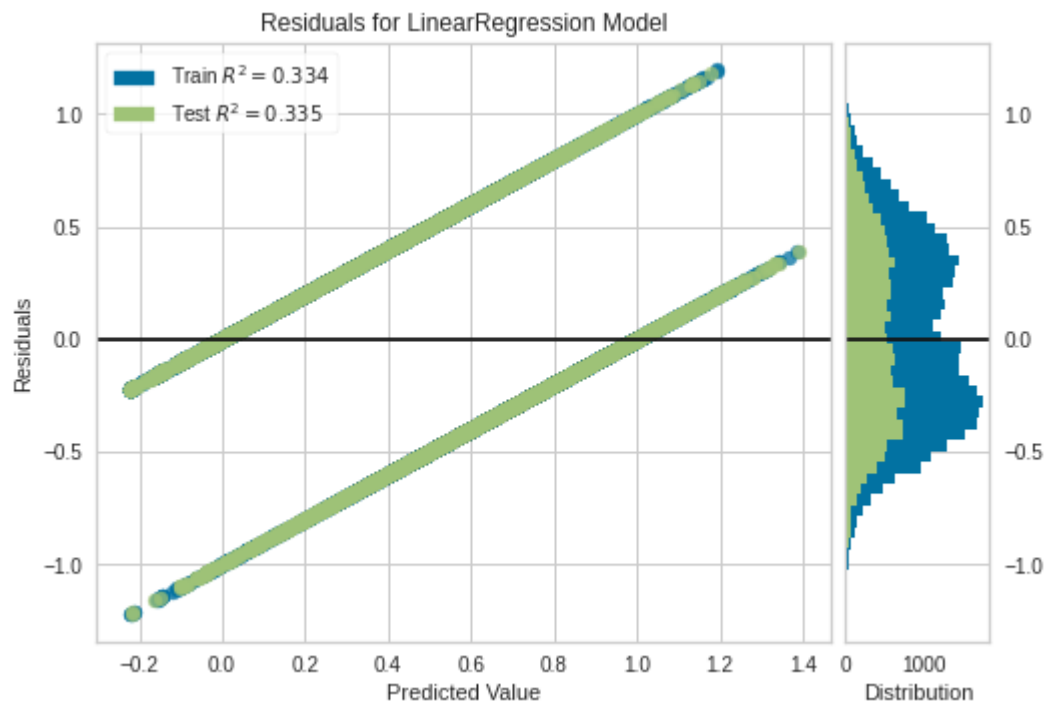
print("Linear Regression Accuracy",score)
```

Linear Regression Accuracy 0.3319457969522649

```
In [30]: from yellowbrick.regressor import ResidualsPlot
visual = ResidualsPlot(lin_reg)
visual.fit(X_train,y_train)
visual.score(X_test,y_test)
visual.poof()
```

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:451: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names

"X does not have valid feature names, but"




```
Out[30]: <AxesSubplot:title={'center':'Residuals for LinearRegression Model'}, xlabel='Predicted Value', ylabel='Residuals'>
```

## Decision Tree

```
In [63]: dt_clf = DecisionTreeClassifier(max_depth=4)
dt_clf.fit(X_train, y_train)
dt_pred = dt_clf.predict(X_test)
dt_acc = metrics.accuracy_score(y_test, dt_pred)
print("Decision Tree Accuracy",100*dt_acc)
```

Decision Tree Accuracy 72.94977168949772

```
In [53]: con_matrix = confusion_matrix(y_test,dt_pred)
sns.heatmap(con_matrix,annot=True,fmt="d")
plt.show()
```



```
In [ ]:
```

```
In [55]: size = 0.1
accuracy_list=[]
dataset_ratio = []
for i in range(1,10):
    print('Test set size: ',(size))
    index = int((i/10)*len(new_df['HeartDisease']))
    X_train, X_test, y_train, y_test = train_test_split(features, target, test_s
size = (i/10), random_state=100)
    dt_clf1 = DecisionTreeClassifier()
    dt_clf1.fit(X_train, y_train)
    dt_pred = dt_clf1.predict(X_test)
    dt_acc1 = metrics.accuracy_score(y_test, dt_pred)
    accuracy_list.append(100*dt_acc1)
    print("Accuracy: ",100*dt_acc1)
    size = round(size, 1)
    dataset_ratio.append(str(i/10)+'%')
    size = round((size+0.1), 1)
```


```
Test set size: 0.1
Accuracy: 68.18264840182648
Test set size: 0.2
Accuracy: 68.80365296803653
Test set size: 0.3
Accuracy: 68.11373599610326
Test set size: 0.4
Accuracy: 67.6743230284488
Test set size: 0.5
Accuracy: 67.58849961641033
Test set size: 0.6
Accuracy: 67.07866536775451
Test set size: 0.7
Accuracy: 67.15027529160035
Test set size: 0.8
Accuracy: 66.76713016873302
Test set size: 0.9
Accuracy: 66.39673648319533
```

do sampling



```
In [57]: g = sns.barplot(x = dataset_ratio, y=accuracy_list)
g.set_ylim(65,70)
```

Out[57]: (65.0, 70.0)



For our conclusion, this dataset does not suitable for classification in Heart Disease cause data overlapping, unbalanced data, low correlation and a lot of outliers. One single feature might not say much but combinations might say more than that! I will perform this in another notebook

```
In [45]: from sklearn import tree
```

```
In [46]: fig = plt.figure(figsize=(25,25))
v = tree.plot_tree(dt_clf,
                  feature_names=features.columns,
                  class_names=['Yes', 'No'],
                  filled=True)
```

